

SlidesGen-Bench: Evaluating Slides Generation via Computational and Quantitative Metrics

Verified reference summary • Year: 2026 • Source bibliography: docs/REFERENCES.md

What the paper contributes

SlidesGen-Bench evaluates slide-generation systems across reproducible axes such as content, aesthetics, and editability. Its value is broader than any single metric: it argues that slide generation must be judged as a multi-dimensional artifact, not just as text-to-image output.

Method / evaluation stance

The benchmark setting naturally uses rendered outputs because systems may emit different formats. Rendering creates a common comparison surface for human preference and visual scoring, which is useful for research fairness.

How slidekit uses it

slidekit takes a complementary route. Because it controls the source format, it can preserve editability and semantics directly in YAML. A title is still a title, a theme color is still a theme role, and a component key still indicates the intended design pattern.

YAML vs HTML takeaway

For production, this makes YAML preferable: editability is built in, not recovered after rendering. For cross-system research, HTML or screenshots remain attractive because they normalize outputs from many generators.

Bottom line

The paper supports a pragmatic split: use research-style rendered artifacts when the question depends on final pixels or human preference collection, but keep slidekit production decks in a constrained YAML IR so layout, linting, and agent fixes remain deterministic, inspectable, and cheap.

Primary slidekit mapping	balance / alignment / whitespace / overlap / contrast / richness / consistency, depending on the paper
Operational risk	HTML metrics can drift with rendering environment; YAML metrics can miss pixel-only effects. Use each where it is strongest.