

Seeing Like a Designer Without One: Unsupervised Slide Quality Assessment via Designer Cue Augmentation

Verified reference summary • Year: 2025 • Source bibliography: docs/REFERENCES.md

What the paper contributes

This work uses expert-inspired visual-design cues for unsupervised slide quality assessment. The cue set includes whitespace, colorfulness, edge density, brightness contrast, text density, color harmony, and layout balance.

Method / evaluation stance

The paper is useful because it separates cues that slidekit can compute from source from cues that require pixels. Whitespace, text density, layout balance, and color roles are available in the IR and ResolvedDeck. Edge density, image texture, and brightness statistics are raster properties.

How slidekit uses it

slidekit therefore adopts the geometry/theme subset and explicitly avoids claiming coverage of pixel-only cues. This is a strength: the system is honest about what a no-render pipeline can and cannot prove.

YAML vs HTML takeaway

HTML/render-first research can see more of the final artifact and is better for these pixel cues. YAML/source-first generation is better for fast, deterministic fixes and for preventing layout defects before a deck is emitted.

Bottom line

The paper supports a pragmatic split: use research-style rendered artifacts when the question depends on final pixels or human preference collection, but keep slidekit production decks in a constrained YAML IR so layout, linting, and agent fixes remain deterministic, inspectable, and cheap.

Primary slidekit mapping	balance / alignment / whitespace / overlap / contrast / richness / consistency, depending on the paper
Operational risk	HTML metrics can drift with rendering environment; YAML metrics can miss pixel-only effects. Use each where it is strongest.