

PPTAgent: Generating and Evaluating Presentations Beyond Text-to-Slides

Verified reference summary • Year: 2025 • Source bibliography: docs/REFERENCES.md

What the paper contributes

PPTAgent and PPTEval broaden evaluation beyond isolated slide text. They treat presentations as structured artifacts with content, design, and coherence dimensions.

Method / evaluation stance

For slidekit, the important idea is the separation between per-slide design quality and deck-level consistency. A single slide can be clean while the deck still feels incoherent because margins, colors, or structure vary too much.

How slidekit uses it

slidekit reflects this by scoring individual resolved slides and applying deck-level consistency logic. YAML helps because themes, component choices, and page-number settings are explicit across the whole deck.

YAML vs HTML takeaway

HTML is useful for holistic visual evaluation, but coherence is easier to preserve when the source format contains a shared theme and typed components. In other words, HTML can show inconsistency; YAML can prevent much of it.

Bottom line

The paper supports a pragmatic split: use research-style rendered artifacts when the question depends on final pixels or human preference collection, but keep slidekit production decks in a constrained YAML IR so layout, linting, and agent fixes remain deterministic, inspectable, and cheap.

Primary slidekit mapping	balance / alignment / whitespace / overlap / contrast / richness / consistency, depending on the paper
Operational risk	HTML metrics can drift with rendering environment; YAML metrics can miss pixel-only effects. Use each where it is strongest.