

Evaluation Metrics for Automated Typographic Poster Generation

Verified reference summary • Year: 2024 • Source bibliography: docs/REFERENCES.md

What the paper contributes

This paper organizes generated-poster evaluation into legibility, aesthetics, and semantics. Posters are not slide decks, but they share core constraints: text must be visible, hierarchy must be clear, and visual structure must support meaning.

Method / evaluation stance

The strongest transfer to slidekit is treating legibility as a constraint rather than a preference. If text is too small, low contrast, or hidden by layout problems, no amount of aesthetic balance can rescue the slide.

How slidekit uses it

slidekit encodes this in safe fonts, minimum body sizes, contrast scoring, and overflow linting. YAML makes these constraints explicit before rendering; the linter can return concrete repair suggestions.

YAML vs HTML takeaway

HTML/CSS offers strong typographic control, but arbitrary CSS can also create hidden overflow and inconsistent font behavior. A constrained YAML IR trades some freedom for reliable legibility guarantees.

Bottom line

The paper supports a pragmatic split: use research-style rendered artifacts when the question depends on final pixels or human preference collection, but keep slidekit production decks in a constrained YAML IR so layout, linting, and agent fixes remain deterministic, inspectable, and cheap.

Primary slidekit mapping	balance / alignment / whitespace / overlap / contrast / richness / consistency, depending on the paper
Operational risk	HTML metrics can drift with rendering environment; YAML metrics can miss pixel-only effects. Use each where it is strongest.