

Learning Layouts for Single-Page Graphic Designs

Verified reference summary • Year: 2014 • Source bibliography: docs/REFERENCES.md

What the paper contributes

O'Donovan, Agarwala, and Hertzmann learn layout preferences for single-page graphic designs using energy terms such as alignment, overlap, whitespace, importance, and balance.

Method / evaluation stance

The paper matters to slidekit because it bridges hand-designed metrics and learned weights. The terms themselves are interpretable, but their combination can be tuned from examples.

How slidekit uses it

slidekit currently uses transparent heuristic weights and avoids human-correlation claims. This paper points to the next step: keep the geometry-based terms, but learn or calibrate weights from labelled slide pairs when data exists.

YAML vs HTML takeaway

HTML/rendered examples are helpful for collecting training preferences. YAML is helpful for deploying the resulting constraints because every element and theme role remains editable and traceable.

Bottom line

The paper supports a pragmatic split: use research-style rendered artifacts when the question depends on final pixels or human preference collection, but keep slidekit production decks in a constrained YAML IR so layout, linting, and agent fixes remain deterministic, inspectable, and cheap.

Primary slidekit mapping	balance / alignment / whitespace / overlap / contrast / richness / consistency, depending on the paper
Operational risk	HTML metrics can drift with rendering environment; YAML metrics can miss pixel-only effects. Use each where it is strongest.