

# Quarterly Business Review

# What Drives Us

## **Excellence**

High standards in every deliverable we ship.

## **Integrity**

Transparent and honest in all we do.

## **Innovation**

Investing in ideas that shape tomorrow.

# By the Numbers

**127%**

Revenue Target

**94%**

Retention Rate

**48 NPS**

Net Promoter

# Side-by-Side Evaluation

## Legacy Platform

Deploy time: 4 hours

Manual scaling

On-premise; high cost

## Modern Stack

Deploy time: 8 min

Auto-scales fast

Cloud-native; low cost

# Key Milestones

**Q1**

## **Foundation**

Cloud platform migration done.

**Q2**

## **API Launch**

REST API live with 99.99% SLA.

**Q3**

## **AI Beta**


AI features shipped to 500 partners.

**Q4**

## **Scale Out**

APAC regions live; global coverage.

# Built in the Field



Built with customers, not just for them.

60% faster rollout; 50% fewer tickets.

# What We Offer

## Insights

Live data  
and AI  
analytics.

## Security

Zero-trust,  
SOC 2  
ready.

## Speed

Sub-100ms  
p99 latency.

## APIs

300+ open  
connectors.

# **Annual Strategy Review**

Executive Leadership — January 2026

# 2025 Results at a Glance

**127%**

Revenue Target

\$3.8B total

**94%**

Retention Rate

**48 NPS**

Net Promoter

+12 pts YoY

# What Drove Growth

## Enterprise & Products

Enterprise up 42%

APAC beat forecast 31%

New products: 18%

## Core & Channel

SMB churn down

Services attach: 34→51%

Partner channel: 28%

# Strategic Priorities 2026

**Q1**

**Foundation**

Merge  
legacy  
systems;  
unified data  
layer.

**Q2**

**AI Assist**

AI  
workflows  
GA for all  
customers.

**Q3**

**Expansion**

Fortune  
500 sales  
motion  
launched.

**Q4**

**Open Beta**

ISV  
program;  
target 50  
partners.

# Investment Areas

## **Build**

Reliability  
and dev  
tooling.

## **Growth**

Enterprise  
sales and  
channels.

## **Global**

APAC and  
EMEA  
expansion.

## **Analytics**

Customer  
and usage  
insights.

# Component Gallery

All eight components in one deck

# Icon Text Rows

## Speed

Deploy in minutes.

## Security

Safe by default.

## Scale

Grows with you.

# Stat Callout

**99.9%**

Uptime

**< 50ms**

Latency

p99 globally

**500+**

Integrations

# Comparison Columns

## Before

Manual; error-prone

Weekly reports

## After

Automated; reliable

Real-time data

# Timeline

## Month 1

### Kickoff

Setup and  
team  
training.

## Month 3

### Go Live

Full  
production  
rollout.

## Month 6

### Growth

Tuning and  
expansion.

# Image Half Bleed

Full-bleed image  
with text side.



# Card Grid

## Feature A

Fast and  
reliable.

## Feature B

Secure  
by  
design.

## Feature C

Scales  
with you.

## Feature D

Pay as  
you go.



01

# Where We Stand Today

# What We Will Cover

---

- 01 Performance against targets
- 02 Customer and market signals
- 03 Operational health and risks
- 04 Roadmap for next quarter
- 05 Decisions we need today

3.2M

---

Active customers worldwide



**"We stopped buying tools and started buying outcomes."**

- VP of Engineering, enterprise customer

# Net Revenue Retention

The percentage of recurring revenue retained from existing customers over a period, including expansion and after churn.



**What would we build if we  
trusted our customers  
completely?**

# What Changed This Quarter

Activation time dropped by half

Support tickets fell sharply

Expansion revenue outpaced new sales



hello.py

```
# The classic, end to end  
def main():  
    print("Hello, world!")
```

main()

```
$ python hello.py  
Hello, world!
```

# Three Steps to Onboard

## **01** Connect your data

Link a source in minutes with a guided wizard.

## **02** Pick a template

Start from a proven layout, then tailor it.

## **03** Share the result

Publish to your team and gather feedback fast.

# Onboarding, Reworked

## Before

---

Manual account setup

Three-day approval wait

Email-only support

## After

---

One-click provisioning

Same-day approval

In-product live chat

# Cloud or On-Prem

**3 wks**

Cloud time-to-launch

vs

**6 mo**

On-prem time-to-launch

# Key Metrics This Quarter

**98.7%**

Uptime

**1,240**

New accounts

**4.6/5**

Support rating

**\$2.1M**

Net new revenue


**12 days**

Avg. time to  
value

**31%**

Expansion rate

# Monthly Active Teams



**Activations  
nearly  
quadrupled  
once self-serve  
onboarding  
shipped in  
February.**

# Pricing Tiers

| Plan       | Seats     | Price     |
|------------|-----------|-----------|
| Starter    | Up to 5   | \$0       |
| Team       | Up to 50  | \$12/seat |
| Business   | Up to 250 | \$20/seat |
| Enterprise | Unlimited | Custom    |

# Performance vs Last Year

**\$8.4M**

Annual recurring  
revenue

**+24%**

**2.1%**

Monthly churn

**-0.8 pts**

**63**

Net promoter  
score

**+11**

# Onboarding Flow

**01**

**Sign up**

Create an  
account

**02**

**Connect**

Link your  
data

**03**

**Configure**

Pick your  
metrics

**04**

**Launch**

Go live in  
minutes

# 2026 Plan

Q1

Ship API

Beta launch

Q2

Mobile app

Integrations

Q3

Enterprise tier

SSO

# Sales Pipeline

Visitors · 50,000

Signups · 8,000

Trials · 2,400

Customers · 600

# Hierarchy of Needs

**Vision**

**Strategy**

**Goals**

**Tactics**

# SWOT Analysis

## Strengths

---

A strong brand with loyal users

## Opportunities

---

New markets through partnerships

## Weaknesses

---

Thin margins and a small team

## Threats

---

New entrants driving price pressure

# Evaluation Matrix

**Platform A**

**Platform B**

**Platform C**

**Security**

**Strong**

Medium

Strong

**Workflow**

Medium

Strong

Medium

**Support**

**Strong**

Strong

Medium

# Job Search Funnel (2021)



# Final folder shape

```
simple-eda-project/  
  src/  
    simple_eda/  
      __init__.py  
      core.py  
  README.md  
  pyproject.toml  
  LICENSE
```

# Agents in AI

How software learned to act, not just answer

# What Makes Software an Agent

## **Perceive**

Read its environment.

## **Reason**

Plan a course of action.

## **Act**

Use tools to change things.

# Assistant vs Agent

## Assistant

Answers one prompt

Waits for you

Has no tools

## Agent

Pursues a goal

Acts on its own

Calls tools

# Core Capabilities

## Planning

Break goals into steps.

## Tool Use

Call APIs and code.

## Memory

Recall past context.

## Reflection

Critique and retry.

# Why Now

**2022**

LLMs Arrive

tool use emerges

**10x**

Faster

early reports

**1M+**

Developers

building agents

# From Rules to Agents

**1980s**

**Rules**

Hand-coded  
logic.

**2010s**

**ML**

Learned patterns.

**Today**

**Agents**

Models that act.

# Common Agent Patterns

## ReAct

Reason,  
then act.

## Tool Use

Call  
functions.

## RAG

Retrieve  
facts first.

## Multi-Agent

Specialists  
team up.

# Multi-Agent Systems

Many narrow agents beat one generalist.

Each owns a clear role.

A coordinator routes the work.

They hand off context.

# Promise vs Pitfalls

## Promise

Automate real work

Scale expertise

Run all day

## Pitfalls

Errors compound

Hard to predict

Needs guardrails

# Takeaways

Software that acts, with human oversight

# Claude Code

A field guide for data scientists

# What We Will Cover

---

- 01 What Claude Code is
- 02 How it differs from chat
- 03 Getting set up
- 04 Everyday data science workflows
- 05 Habits that make it shine



# Claude Code

A terminal-based AI agent that reads your project, writes and runs code, inspects the results, and iterates until the job is done.

# Chat vs Agent

## Chat

You ferry code around

vs

## Agent

It works in your repo

# What It Can Do in Your Project

## Explore

Read files, schemas, and notebooks.

## Execute

Run scripts, queries, and tests.

## Iterate

Check outputs and fix its own errors.

# Set Up in Three Steps

## 01 Install

`npm install -g @anthropic-ai/claude-code`

## 02 Open your project

cd into the repo and run claude.

## 03 Ask for an analysis

Describe the outcome; it plans the code.



a first session

```
$ cd sales-analysis && claude
```

```
> Profile sales.csv: nulls, dtypes,  
  outliers. Save it to profile.md.
```

```
Reading sales.csv (1.2M rows)...
```

```
Running profiling script...
```

```
Wrote profile.md: 3 findings.
```

# Everyday Workflows

## EDA

Profile new data in minutes.

## Cleaning

Fix types, nulls, duplicates.

## Modeling

Fast first baselines.

## Reporting

Plots and write-ups.

# How It Works on a Task

**01**

## **Explore**

Scans the data and code.

**02**

## **Plan**

Proposes an approach.

**03**

## **Execute**


Writes and runs the code.

**04**

## **Verify**

Checks results, then iterates.

# One Prompt, a Whole Pipeline



# A Typical Day, Reworked

## Before

Boilerplate pandas by hand

Stack Overflow detours

Stale, undocumented  
notebooks

## With Claude Code

Describe the outcome  
instead

Errors debugged in place

Analysis committed with  
notes

# Commands Worth Knowing

## Command

## What it does

---

`claude`

Start a session

`/init`

Write CLAUDE.md

`/compact`

Trim the context

`claude -p`

One-shot mode

# Habits That Make It Shine

Keep a CLAUDE.md with data caveats.

Ask for a plan before a big analysis.

Have it verify outputs, not just run.

Commit small, reviewable steps.



**What would you analyze if  
the boilerplate wrote itself?**

# Start Today

Open a terminal and ask

# **slidekit First Light**

Phase 5 demo: lint-clean, zero screenshots

# How slidekit Works

Author IR YAML once

Layout engine solves EMU rects

Linters validate before python-pptx emits

Round-trip test verifies — zero screenshots

# Calibration Results

**0.076%**

Max Width Error

vs LibreOffice render

**2%**

Error Threshold

hard gate PASS

**122**

Tests Passing

phases 1 through 5

# Design Principles

## **Deterministic**

Same input always produces same output.

## **No screenshots**

Lint errors cite path and exact fix.

## **Agent-ready**

Build and fix loop needs no images.

# Before and After

## Ad-hoc pptx code

Manual EMU math

Visual QA required

No lint feedback

Agent loops on images

## slidekit

Declarative IR YAML

Round-trip verified

Lint errors with fix

Agent loops on text

# **Interface Aesthetics**

Ngo, Teo & Byrne · Information Sciences (2003)



**Is a layout's beauty  
subjective — or can it be  
computed?**



# Interface aesthetics

The visual quality of a layout, written as numbers from where each element sits and how big it is.

# 14

---

## Aesthetic measures

Closed-form functions of element bounding boxes

# The core claim

Layout quality is not purely subjective.

Balance and spacing become numbers.

Each number is a formula over boxes.

Measured, optimised, checked vs people.



01

## **The fourteen measures**

# The measures (1–5)

Balance — even visual weight

Equilibrium — mass near the centre

Symmetry — mirrored about the axes

Sequence — a clear reading flow

Cohesion — aspect ratios agree



# Bounding box

Every measure reads only position and size — x, y, width, height. No pixels.

# The measures (6–10)

Unity — parts read as one

Proportion — pleasing ratios

Simplicity — few alignment points

Density — comfortable coverage

Regularity — elements snap to a grid

# The measures (11–14)

Economy — few sizes, reused

Homogeneity — evenly distributed

Rhythm — regular, repeating change

Order & complexity — structure wins

regularity – Ngo (2003)

# Ngo regularity (alignment term)

$$\text{regularity} = 1 - (\text{n\_vap} + \text{n\_hap}) / (2n)$$

# a tidy 5-bullet column:

#  $n = 5$ ,  $\text{n\_vap} = 1$ ,  $\text{n\_hap} = 5$

$$\text{regularity} = 1 - (1 + 5) / 10 \quad \# 0.40$$

# slidekit: left edges only

$$\text{regularity} = 1 - 1 / 5 \quad \# 0.80$$

# Same clean column, two scores

**0.40**

Ngo x + y regularity

vs

**0.80**

Left-edge alignment

# Why slidekit fits this paper

Every element already has an exact rect.

Measures are arithmetic over geometry.

Deterministic — same deck, same score.

Zero rendering, zero vision.



02

## **From paper to slidekit**

# What slidekit adopted

## Declined

Regularity

Rhythm

Proportion

Economy

Spacing-regularity

## Adopted

Balance

Density

# Adoption decisions

| Measure    | Verdict  | Reason      |
|------------|----------|-------------|
| Balance    | Adopted  | Axis moment |
| Density    | Adopted  | Coverage    |
| Regularity | Declined | Halves cols |
| Proportion | Declined | Ratio noise |

# A measured-aesthetics lineage

1933

**Birkhoff**

A beauty  
ratio

2003

**Ngo**

Fourteen  
measures

2004

**Lok**

Visual  
balance

2026

**Zhang**

Moment  
model

# Honest scope

The journal paper is paywalled.

Public copies reproduce only alignment.

Constants are not copied blindly.

No universal beauty equation is claimed.

# From measure to code

## 01 Read geometry

Take each element's rect.

## 02 Define a formula

Write an open, closed-form score.

## 03 Verify & ship

Check the source, then report.



**"Turn design quality into explicit, auditable measures."**

- Ngo, Teo & Byrne (2003)