

Quarterly Business Review

What Drives Us

Excellence

High standards in every deliverable we ship.

Integrity

Transparent and honest in all we do.

Innovation

Investing in ideas that shape tomorrow.

By the Numbers

127%

Revenue Target

94%

Retention Rate

48 NPS

Net Promoter

Side-by-Side Evaluation

Legacy Platform

Deploy time: 4 hours

Manual scaling

On-premise; high cost

Modern Stack

Deploy time: 8 min

Auto-scales fast

Cloud-native; low cost

Key Milestones

Q1

Foundation

Cloud platform migration done.

Q2

API Launch

REST API live with 99.99% SLA.

Q3

AI Beta


AI features shipped to 500 partners.

Q4

Scale Out

APAC regions live; global coverage.

Built in the Field



Built with customers.

60% faster; 50% fewer tickets

What We Offer

Insights

Live data
and AI
analytics.

Security

Zero-trust,
SOC 2
ready.

Speed

Sub-100ms
p99 latency.

APIs

300+ open
connectors.

Annual Strategy Review

Executive Leadership — January 2026

2025 Results at a Glance

127%

Revenue Target

\$3.8B total

94%

Retention Rate

48 NPS

Net Promoter

+12 pts YoY

What Drove Growth

Enterprise & Products

Enterprise up 42%

APAC beat forecast 31%

New products: 18%

Core & Channel

SMB churn down

Services attach: 34→51%

Partner channel: 28%

Strategic Priorities 2026

Q1

Foundation

Merge legacy systems; unified data layer.

Q2

AI Assist

AI workflows GA for all customers.

Q3

Expansion

Fortune 500 sales motion launched.

Q4

Open Beta

ISV program; target 50 partners.

Investment Areas

Build

Reliability
and dev
tooling.

Growth

Enterprise
sales and
channels.

Global

APAC and
EMEA
expansion.

Analytics

Customer
and usage
insights.

Component Gallery

All eight components in one deck

Icon Text Rows

Speed

Deploy in minutes.

Security

Safe by default.

Scale

Grows with you.

Stat Callout

99.9%

Uptime

< 50ms

Latency

p99 globally

500+

Integrations

Comparison Columns

Before

Manual; error-prone

Weekly reports

After

Automated; reliable

Real-time data

Timeline

Month 1

Kickoff

Setup and
team
training.

Month 3

Go Live

Full
production
rollout.


Month 6

Growth

Tuning and
expansion.

Image Half Bleed

Image with text side.



Card Grid

Feature A

Fast and
reliable.

Feature B

Secure
by
design.

Feature C

Scales
with you.

Feature D

Pay as
you go.



01

Where We Stand Today

What We Will Cover

- 01 Performance against targets
- 02 Customer and market signals
- 03 Operational health and risks
- 04 Roadmap for next quarter
- 05 Decisions we need today

3.2M

Active customers worldwide



"We buy outcomes, not tools."

- VP of Engineering, enterprise customer



Net Revenue Retention

Revenue kept from existing customers.

| What if we trusted users?

What Changed This Quarter

Activation time dropped by half

Support tickets fell sharply

Expansion revenue outpaced new sales



hello.py

```
# The classic, end to end  
def main():  
    print("Hello, world!")
```

```
main()
```

```
$ python hello.py  
Hello, world!
```

Three Steps to Onboard

01 Connect your data

Link a source in minutes with a guided wizard.

02 Pick a template

Start from a proven layout, then tailor it.

03 Share the result

Publish to your team and gather feedback fast.

Onboarding, Reworked

Before

Manual account setup

Three-day approval wait

Email-only support

After

One-click provisioning

Same-day approval

In-product live chat

Cloud or On-Prem

3 wks

Cloud time-to-launch

vs

6 mo

On-prem time-to-launch

Key Metrics This Quarter

98.7%

Uptime

1,240

New accounts

4.6/5

Support rating

\$2.1M

Net new revenue


12 days

Time to value

31%

Expansion rate

Monthly Active Teams



Pricing Tiers

Plan	Seats	Price
Starter	Up to 5	\$0
Team	Up to 50	\$12/seat
Business	Up to 250	\$20/seat
Enterprise	Unlimited	Custom

Performance vs Last Year

\$8.4M

Revenue

+24%

2.1%

Monthly churn

-0.8 pts

63

Promoter score

+11

Onboarding Flow

01

Sign up

Create an
account

02

Connect

Link your
data

03

Configure

Pick your
metrics

04

Launch

Go live in
minutes

2026 Plan

Q1

Ship API

Beta launch

Q2

Mobile app

Integrations

Q3

Enterprise tier

SSO

Sales Pipeline

Visitors · 50,000

Signups · 8,000

Trials · 2,400

Customers · 600

Hierarchy of Needs

Vision

Strategy

Goals

Tactics

SWOT Analysis

Strengths

Strong brand, loyal users

Opportunities

New markets via partners

Weaknesses

Thin margins, small team

Threats

New entrants cut prices

Evaluation Matrix

	Platform A	Platform B	Platform C
Security	Strong	Medium	Strong
Workflow	Medium	Strong	Medium
Support	Strong	Strong	Medium

Job Search Funnel (2021)



Final folder shape

```
simple-eda-project/  
  ■■ src/  
    ■ ■■ simple_eda/  
      ■ ■■ __init__.py  
      ■ ■■ core.py  
    ■■ README.md  
    ■■ pyproject.toml  
    ■■ LICENSE
```

Agents in AI

How software learned to act, not just answer

What Makes Software an Agent

Perceive

Read its environment.

Reason

Plan a course of action.

Act

Use tools to change things.

Assistant vs Agent

Assistant

Answers one prompt

Waits for you

Has no tools

Agent

Pursues a goal

Acts on its own

Calls tools

Core Capabilities

Planning

Break goals into steps.

Tool Use

Call APIs and code.

Memory

Recall past context.

Reflection

Critique and retry.

Why Now

2022

LLMs Arrive

tool use emerges

10x

Faster

early reports

1M+

Developers

building agents

From Rules to Agents

1980s

Rules

Hand-coded.

2010s

ML

Learned patterns.

Today

Agents

Models that act.

Common Agent Patterns

ReAct

Reason,
then act.

Tool Use

Call
functions.

RAG

Retrieve
facts first.

Multi-Agent

Specialists
team up.

Multi-Agent Systems

Many narrow agents beat one generalist.

Each owns a clear role.

A coordinator routes the work.

They hand off context.

Promise vs Pitfalls

Promise

Automate real work

Scale expertise

Run all day

Pitfalls

Errors compound

Hard to predict

Needs guardrails

Takeaways

Software that acts, with human oversight

Claude Code



Claude Code

A terminal AI agent that writes and runs your code.

Chat vs Agent

Chat

You ferry code around

vs

Agent

It works in your repo

What It Can Do in Your Project

Explore

Read files, schemas, and notebooks.

Execute

Run scripts, queries, and tests.

Iterate

Check outputs and fix its own errors.



Terminal Use

```
$ cd sales-analysis && claude
```

```
> Profile sales.csv: nulls, dtypes,  
outliers. Save it to profile.md.
```

```
Reading sales.csv (1.2M rows)...
```

```
Running profiling script...
```

```
Wrote profile.md: 3 findings.
```

Everyday Workflows

EDA

Profile new data in minutes.

Cleaning

Fix types, nulls, duplicates.

Modeling

Fast first baselines.

Reporting

Plots and write-ups.

How It Works on a Task

01

Explore

Scans the data and code.

02

Plan

Proposes an approach.

03

Execute

Writes and runs the code.

04

Verify

Checks results, then iterates.

A Typical Day, Reworked

Before

Boilerplate pandas by hand

Stack Overflow detours

Stale notebooks

With Claude Code

Describe the outcome

Errors debugged in place

Committed with notes

Commands Worth Knowing

Command	What it does
<code>claude</code>	Start a session
<code>/init</code>	Write CLAUDE.md
<code>/compact</code>	Trim the context
<code>claude -p</code>	One-shot mode

Two Commands, Up Close

/init

Scans your project, writes a starter CLAUDE.md.

claude -p

One-shot: run a prompt, print the answer, exit.

Habits That Make It Shine

Keep a CLAUDE.md with data caveats.

Ask for a plan before a big analysis.

Have it verify outputs, not just run.

Commit small, reviewable steps.



01

One Claude, Many Surfaces

Same Engine, Everywhere

One engine: terminal, IDE, desktop, web.

Your CLAUDE.md and MCP servers come too.

Start on the laptop, finish on your phone.

Three Ways to Run It

	Terminal	Web	Desktop
Runs on	Your machine	The cloud	Your machine
Setup	Install CLI	None	Install app

On the Internet

Cloud-run

Runs at `claude.ai/code` on Anthropic's cloud.

Isolated

Each task gets its own sandbox VM.

Persistent

Keeps running after you close the tab.

Claude Desktop

A standalone app, no terminal needed.

Review diffs, run sessions side by side.

Schedule repeats and start cloud runs.



02

Drive It From Your Phone

Connect GitHub

01 Sign in

Open `claude.ai/code` and log in.

02 Link GitHub

Grant access to the repos you choose.

03 Pick a repo

Claude clones it into a cloud sandbox.

Start a Task Anywhere

01 Describe it

Type the task in the Claude mobile app.

02 Let it run

Work continues in the background.

03 Review

Approve the branch Claude pushes.

Remote Control

Step away and work from your phone.

Continue a live session from any browser.



Move Between Surfaces

```
$ claude --cloud
```

Runs this task **in** the cloud.

> Pick it up later, anywhere:

```
$ claude --teleport
```

One session follows you across
laptop, phone, **and** browser.



03

MCP: Plug In Your Tools



MCP

An open standard linking AI tools to your data.

What MCP Unlocks

Read design docs from Google Drive.

Update Jira, pull data from Slack.

Query your own database and APIs.

Without MCP vs With MCP

Without MCP

Paste data in by hand

Screenshot dashboards

Describe your schema

With MCP

Claude queries the DB

Reads the ticket itself

Always current data

Servers You Can (normally) Connect to

Server	What Claude can do
Drive	Read design docs
Jira	Update tickets
Slack	Pull conversations
Postgres	Query your data



04

Customize Claude



CLAUDE.md

A project file Claude reads every session.

What Goes in CLAUDE.md

Coding standards and libraries.

Architecture and key commands.

Review checklists and data caveats.

Example CLAUDE.md

sales-analysis project

- Use pandas, **not** polars.
- Raw data lives **in** data/.
- Never commit the CSVs.
- Tests: pytest -q
- Plots: matplotlib only.

CLAUDE.md Pitfalls

Too long crowds the context.

Stale rules Claude still obeys.

Vague advice it cannot act on.

Secrets do not belong here.



Skills

Reusable instruction folders, loaded on demand.

CLAUDE.md vs Skills.md

CLAUDE.md

Always in context

Project-wide rules

One file at the root

Skills.md

Loaded on demand

A packaged workflow

Shared across projects



05

More From Claude



Cowork

Claude Code in Claude Desktop, for non-code work.

What Cowork Does

Reads and edits files you choose.

Works across Slack and Google Drive.

Finishes multi-step tasks, with citations.

Artifacts

Shareable code, docs and apps built beside chat.

Memory > Chat Search

Memory

Learns your context

Persists across chats

Separate per project

Chat search

Looks up old chats

Only when you ask

One lookup at a time

| Build it from our phones.

slidekit First Light

Phase 5 demo: lint-clean, zero screenshots

How slidekit Works

Author IR YAML once

Layout engine solves EMU rects

Linters validate before python-pptx emits

Round-trip test verifies — zero screenshots

Calibration Results

0.076%

Max Width Error

vs LibreOffice render

2%

Error Threshold

hard gate PASS

122

Tests Passing

phases 1 through 5

Design Principles

Deterministic

Same input always produces same output.

No screenshots

Lint errors cite path and exact fix.

Agent-ready

Build and fix loop needs no images.

Before and After

Ad-hoc pptx code

Manual EMU math

Visual QA required

No lint feedback

Agent loops on images

slidekit

Declarative IR YAML

Round-trip verified

Lint errors with fix

Agent loops on text

Build Your First Library

The path to PyPI

01

Function

pandas
code

02

Package

one folder

03

Build

make a
wheel

04

Publish

pip install

Today's goals

01 Package

Create a real, importable Python package.

02 Publish

Upload it to TestPyPI, then PyPI.

03 Install

`pip install` your own package.

| Build tiny. Publish safely.



What we're building

```
import pandas as pd  
import simple_eda as eda
```

```
df = pd.read_csv("data.csv")  
eda.summarize(df)  
eda.missing(df)
```

What we build now — and skip

Build now

One dependency: pandas

Plain Python functions

A clear README

Skip for now

CI/CD

Testing

Many dependencies



01

What is a Python library?



Library

Reusable code you install once and import anywhere.

Script vs library

Script

`python analyze.py`

Less convenient

Limited scope

Library

`import simple_eda`

More convenient

Broad scope

A tiny library is still real

No web app required.

No dashboard required.

Just useful, installable, importable.

The five bare minimums

A unique package name

A project folder

An `__init__.py` file

Your functions

Metadata — `pyproject.toml`



02

Design the smallest API



API

The functions and classes users call.



MVP

The smallest version that is genuinely useful.



the simplest API

```
import simple_eda as eda  
eda.summarize(df)
```

Pick boring function names

Prefer (boring)

summarize

missing

numeric_columns

Avoid (clever)

hyper_profile_engine

magic_diagnose

mega_summary

MVP Input: Pandas dataframe only

Support pandas DataFrames first.

Not Polars, Spark, Dask, or Arrow yet.

Fewer choices means fewer bugs.

MVP Output: Plain objects only

Dictionaries & lists

Simple built-in containers.

Ints & strings

Plain scalar values.

pandas Series

No custom classes to learn.

The 0.1.0 mindset

First working version, not perfect.

Installable and understandable.

Momentum, not completeness.



03

Create the package folder

Final folder shape

```
simple-eda-project/  
  ■■ src/  
    ■ ■■ simple_eda/  
      ■ ■■ __init__.py  
      ■ ■■ core.py  
  ■■ README.md  
  ■■ pyproject.toml  
  ■■ LICENSE
```

Install name vs import name

Install name

simple-eda-you

hyphens

used by pip & PyPI

Import name

simple_eda

underscores

used in Python code

Why the names differ

import needs a valid Python name.

Python names cannot contain hyphens.

So the import name uses underscores.

PyPI treats - and _ as the same name.



`__init__.py`

Marks a folder as an importable package.

Why the src/ layout?

Blocks accidental root imports.

Matches modern Python projects.

Makes packaging behavior clear.

Common folder mistakes

pyproject.toml in the wrong folder.

Forgetting `__init__.py`.

Folder named `simple-eda`, not `simple_eda`.

Installing from the wrong directory.

What each mistake causes

Mistake	What happens
No <code>__init__.py</code>	import cannot find it
Hyphen in folder	invalid module name
pyproject misplaced	build finds nothing
Wrong directory	ModuleNotFoundError



04

Write the EDA functions



```
def summarize(df):  
    return {  
        "rows": len(df),  
        "columns": len(df.columns),  
        "names": list(df.columns),  
    }
```

The core functions

summarize — shape, columns, types

missing — nulls per column

numeric_columns / categorical_columns



missing()

```
def missing(df):  
    counts = df.isna().sum()  
    return counts.sort_values(  
        ascending=False)
```

```
# missing(df) ->  
# age      2  
# name     1
```

numeric & categorical columns

```
def numeric_columns(df):  
    sel = df.select_dtypes(  
        include="number")  
    return list(sel.columns)
```

```
def categorical_columns(df):  
    sel = df.select_dtypes(  
        include=["object", "category"])  
    return list(sel.columns)
```




Try it on a tiny DataFrame

```
import simple_eda as eda
```

```
df = pd.DataFrame({  
    "name": ["Ana", "Bo", None],  
    "age": [25, None, 31],  
})  
print(eda.summarize(df))  
print(eda.missing(df))
```

Function do's and don'ts

Do

Return data

Use plain dicts / Series

Keep signatures small

One job per function

Don't

Change df in place

Print instead of return

Add many optional args

One function for everything

One job vs everything

One job each

Small and focused

Easy to test

Easy to reuse

One big function

Hard to follow

Hard to test

Breaks more often



Signature

A function name plus the parameters it takes.



05

Add pyproject.toml

What pyproject.toml does

Names your package and version.

Lists dependencies.

Tells build tools how to build.



Build tools

Tools that turn source into installable packages.

Minimal pyproject.toml

```
[build-system]
```

```
requires = ["setuptools>=61"]
```

```
build-backend = "setuptools.build_meta"
```

```
[project]
```

```
name = "simple-eda-yourname"
```

```
version = "0.1.0"
```

```
description = "Tiny EDA helper"
```

```
dependencies = ["pandas>=1.5"]
```




TOML

The simple key = value format pyproject uses.

Version numbers

Version	Means
0.1.0	First working version
0.1.1	A small fix
0.2.0	A new feature
1.0.0	Stable and public



MAJOR.MINOR.PATCH

Each of the three numbers is a non-negative integer.

README: install & use

simple-eda-yourname

Tiny EDA helper for pandas.

Install

```
pip install simple-eda-yourname
```

Use

```
import simple_eda as eda  
print(eda.summarize(df))
```

License basics

Says what others may do with your code.

MIT is common for small projects.

Use an official template.

Never invent your own license text.



06

Install locally with pip

Install in editable mode

```
# run from the folder with pyproject.toml  
python -m pip install -e .
```

What -e (editable) means

Links pip to your source folder.

Code edits apply instantly.

No reinstall while developing.

Smoke-test the import

```
>>> import simple_eda as eda  
>>> eda.__all__  
['summarize', 'missing', ...]
```

When the import fails

Symptom

ModuleNotFoundError

Wrong folder

Wrong environment

Fix

Run `pip install -e .`

`cd` to the project root

Use `python -m pip`



07

Build a real package file

Build the artifacts

```
python -m pip install build
```

```
python -m build
```

The dist/ folder

dist/

■■■ simple_eda_yourname-0.1.0.tar.gz

■■■ simple_eda_yourname-0.1.0-...whl

Wheel vs source distribution

Source (.tar.gz)

A source archive

Used as a fallback

Wheel (.whl)

Prebuilt, fast install

What most users get

Keep build output out of git

dist/, build/, *.egg-info are generated.

Add them to .gitignore.

Bump the version before rebuilding.



08

Publish safely

Why TestPyPI first?

A practice package index.

Rehearse upload and install commands.

Mistakes are cheap here.

The upload flow

01 Check

```
python -m twine check dist/*
```

02 TestPyPI

```
twine upload --repository testpypi dist/*
```

03 PyPI

```
twine upload dist/* — only after TestPyPI works.
```

Install from TestPyPI

```
# install from TestPyPI
```

```
pip install \
  -i https://test.pypi.org/simple/ \
  --extra-index-url \
    https://pypi.org/simple/ \
  simple-eda-yourname
```

Publishing safety

Use API tokens, not passwords.

Never commit tokens or secrets.

Fewer deps, less attack surface.

Beware typosquatting.



09

Use, improve & share

What to add — and avoid

Avoid early

Full HTML reports

Plugin systems

Every dataframe backend

AI-generated analysis

Good next steps

`quick_report()`

`top_categories()`

`column_summary()`

`duplicate_rows()`



Add a tiny test

```
def test_numeric_columns():  
    df = pd.DataFrame(  
        {"name": ["A"], "age": [10]})  
    cols = eda.numeric_columns(df)  
    assert cols == ["age"]
```

Final checklist

Package files all exist.

`pip install -e .` works.

`python -m build` works.

`twine check dist/*` passes.

README has install + usage.

0.1.0



Ship your first version

Four functions, importable, published.

From Vague to Precise

Immersion Week SF · Day 1 · Session 1.1

35



Minutes

You leave with one precise problem statement

What This Session Covers

- 01** Problem statements versus solution statements
- 02** Making Claude interrogate a vague ask
- 03** The three-gate scoping test
- 04** Standing up your Claude Project

| What are you building?

Two Ways to Answer That

Vague

An AI thing for resumes

vs

Precise

Flags gaps in 200 files

How We Get There

01

Separate

Problem
apart from
solution

02

Interrogate

Let Claude
attack the
ask

03

Scope

Run the
three-gate
test

04

Commit

Write it into
the Project



01

Problem, Not Solution



Problem statement

A gap between what happens and what should.



Solution statement

A thing you plan to build, assuming the diagnosis.

How to Tell Them Apart

Problem

Describes a gap

Admits several fixes

Names who suffers

Solution

Describes a build

Assumes one fix

Names a tool

What They Said, What They Meant

Said (solution)

Meant (problem)

Build a RAG chatbot

Staff cannot find answers

Fine-tune a model

Output misses our tone

Make a dashboard

Nobody sees the numbers

Use an agent

Review takes four hours

Why Solution-First Hurts

You cannot tell if it worked

The tool choice is already locked

Scope creeps with no boundary

Feedback arrives too late to use

Three Tells You Wrote a Solution

A tool is named

RAG, agent, fine-tune, dashboard

Only one fix fits

No alternative is even thinkable

No one is in it

Nobody is described as stuck



"Names a tool? Not a problem yet."

- The fastest test in the room

Interrogation

Making Claude attack a vague ask until it is precise.

What Makes Claude Interrogate

Give it the adversarial role

Forbid it from proposing solutions

Demand one question at a time

Make it refuse to move on early



The Interrogation Prompt

You are a skeptical product lead.
My idea: "an AI thing for resumes"

Rules:

- Ask ONE question at a time.
- Never propose a solution.
- Reject vague answers.
- Stop when it **is** precise.

The Interrogation Loop

01 Answer plainly

No hedging, no tool names

02 Let it push

Expect the follow-up to hurt

03 Restate

Ask Claude to write it back

Before and After the Interrogation

Before

An AI thing for resumes

Who for? Unclear

Done when? Unclear

After

Recruiters miss skill gaps

For 200 intern applicants

Flagged in one pass

Signs It Is Now Precise

A stranger can restate it

You know who is stuck

You can name a wrong answer

Two people would agree it is solved



02

The Scoping Test

Scoping test

Three gates a problem passes before you commit.

The Three Gates

Measurable outcome

You can say whether it worked

Data available

You have it, today, in hand

Three days

Finishable by Friday noon



Gate 1 · Measurable

State up front what would prove it worked.

Vague Metric, Measurable Metric

Vague

Better resumes

Faster reviews

Higher quality

More useful

Measurable

Flags 8 of 10 known gaps

Under 30 seconds per file

Two reviewers agree it passes

Used unprompted twice



Gate 2 · Data

It exists, you may use it, you have it today.

Where Data Plans Die

Access pending

Someone still has to approve it

Wrong shape

Exists, but not in usable form

Too sensitive

Real records you cannot touch



Gate 3 · Three days

The whole path working, not one part polished.

What Three Days Buys

One path

End to end and rough

vs

One part

Polished but alone

Three Candidate Ideas, Three Gates

	Resumes	Policy	Grader
Measurable	Yes	Yes	Yes
Data now	Yes	No	Yes
3 days	No	Yes	No

The opposite case

Making Claude argue your problem is the wrong one.



Make Claude Argue Against You

Our problem statement:

"<paste it>"

Argue this **is** the WRONG problem
for a three-day build.

Three strongest objections,
hardest first. No alternatives.

Running the Opposition

01 Read all three

Do not answer while reading

02 Rank them

Which would actually sink you

03 Fix or accept

Narrow it, or note the risk

When the Objection Wins

Narrow the users, not the ambition

Shrink the data, keep the path

Stub the hard step for now

Change problems only if a gate fails

Cut Scope, Not Ambition

One user type, not four

One input format, not any

One clean path, no edge cases



03

The Team's Shared Brain



Claude Project

A shared workspace holding instructions and context.

Chat Versus Project

A chat

Context dies at the tab

One person holds it

Re-explained every time

A Project

Context persists

The whole team shares it

Explained once

What a Project Holds

Instructions

Standing rules for every chat

Context

Uploaded docs and data

One place

Everyone works from it

Stand One Up Now

01 Create

Name it for the problem

02 Instruct

Paste the problem statement

03 Share

Add every teammate today

Project Instructions Template

Problem

<the precise statement>

Success looks like

<the measurable outcome>

Out of scope

<what we cut, and why>

What to Upload

Sample data, even if faked

The interrogation transcript

The objections you accepted

What Not to Upload

Real personal records

Anything under an NDA

Credentials or API keys

Team Conventions Worth Setting

One source

Project over
chat

Log it

Decisions
land in
writing

Shared

Everyone
reads it first

Update

Change it
when scope
does

Project Pitfalls

Instructions written once, never updated

Uploads nobody has actually read

One member still working in chat

Done With This Session When

The statement passes three gates

The strongest objection is answered

The Project exists and is shared

What You Leave With

1

Precise problem

3

Gates passed

1

Shared Project

Next: Mapping the Workflow

Break the problem into steps

Draw it as a Claude Artifact

Name what each step needs

| Can a stranger restate it?

Mapping the Workflow

Immersion Week SF · Day 1 · Session 1.2

55



Minutes

You leave with a diagram the team shares

What This Session Covers

- 01** Decomposing the problem into steps
- 02** Turning the map into an Artifact
- 03** Naming what each step needs
- 04** Finding how the riskiest step fails

| From input to output: how?

From 1.1 to Now

Problem

One precise sentence

vs

Workflow

The path to solve it

How We Get There

01

Break

Split it into
steps

02

Draw

Make it an
Artifact

03

Annotate

Name what
each needs

04

Stress

Find how it
breaks



01

Decompose Into Steps



Workflow

Ordered steps from raw input to delivered output.



Step

One change, with a named input and output.

Every Step Has Three Parts

Input

What arrives, and its shape

Transform

The one thing this step does

Output

What leaves, ready for next

How to Decompose

01 Name the output

What the user finally gets

02 Work backwards

What must exist before it

03 Stop at raw input

What you actually receive

Start From the Output

Forward drifts toward what is easy

Backwards keeps every step earning

You spot the missing middle sooner

Vague Step, Concrete Step

Vague

Concrete

Process the data

Parse CSV to records

Analyze it

Score each row 0 to 1

Use AI

Claude drafts a summary

Show results

Render a ranked table

The Granularity Test

One person could own it

You could test it alone

It fits in a morning

Too Coarse, Too Fine

Too coarse

Build the pipeline

Hides the hard part

Cannot be assigned

Too fine

Open the file

Twenty trivial boxes

Noise, not a map



Prompt to Decompose

Our problem:

"<paste the 1.1 statement>"

Work BACKWARDS from the output.

List the steps, each with its
input and output.

Stop at the raw input we have.

5-9 steps. No tools yet.

A Worked Example

01

Intake

Resume
files arrive

02

Parse

Text out of
each file

03

Score

Match
against the
role

04

Report

Ranked list
with gaps

Common Decomposition Mistakes

Starting after the messy part

One step doing three things

No step for getting the data

The output nobody defined

Linear or Branching

Linear

One path start to end

Easy to demo

Pick this for Day 2

Branching

Cases split the path

Harder to finish

Note it, defer it



The happy path

The one route where nothing fails.

Mark the Happy Path

One line through the map

Every branch labelled later

Edge cases listed, not built



"A gap in the map is a gap in Day 2."

- Why the middle matters



02

Make It an Artifact



Claude Artifact

Standalone content Claude builds, shared by link.

Described or Drawn

Described

Lives in a transcript

Each person pictures it

Re-explained each time

Drawn

One picture, one meaning

Disagreement shows fast

Sent as a link

What an Artifact Gives You

See

The whole path at once

Revise

Change it in seconds

Share

One link, whole team

Build the Diagram

01 Paste the steps

Your list from section one

02 Ask for a diagram

Boxes and arrows, top down

03 Fix it out loud

Revise until it matches



Prompt for the Diagram

Here are our workflow steps:
<paste the list>

Build an Artifact: a flow diagram, top to bottom.

One box per step.

Label every arrow **with** what passes between the steps.

What the Diagram Must Show

Raw input, where it enters

One box per step

What passes on each arrow

The delivered output

Diagram Conventions

Top to bottom reads fastest

One verb per box

Keep labels under five words

Revise It in Front of Everyone

Ask for the change out loud

Re-read the new version together

Stop when nobody objects

Artifact Pitfalls

A pretty diagram nobody rechecks

Boxes with no arrow labels

The link never actually shared

Where the Map Lives

Chat

Scrolls away by noon

vs

Project

Beside the problem



03

Name What Each Step Needs

Four Kinds of Need

Data

Something
must be
supplied

Tool

Code, API,
or model
call

Routine

It runs on a
schedule

Human

A person
decides

Annotate Every Box

One need per step, minimum

Say which of the four it is

Unknowns get marked, not skipped

Claude Routine

A prompt Claude runs on a schedule, unprompted.

When a Step Is a Routine

It repeats on a clock

Nobody should have to remember

Late is worse than imperfect

A Step as a Routine

Every weekday at 8:00

1. Read new files `in /intake`
2. Score each against the role
3. Append to the ranked sheet
4. Flag anything scoring low



Human decision

A step where a person judges and the system waits.

Keep a Human When

Being wrong is expensive

The call needs context you lack

Someone must be accountable

Step to Need

Step	What it needs
Intake files	Data plus a routine
Parse text	A tool
Score match	A model call
Approve list	A human decision



The riskiest step

The one whose failure wastes the most time.

Find It

01 Mark unknowns

Boxes you cannot describe

02 Ask who blocks

Which stalls everything after

03 Pick one

Only one gets this attention



Prompt for Failure Modes

Our riskiest step:

"<describe the step>"

Give the THREE most likely ways
this fails **in** practice.

Most likely first.

For each: the early warning sign.

No fixes yet.

Where Steps Usually Fail

Input

Arrives wrong or not at all

Transform

Works, but not on real data

Access

The key, quota, or approval

What to Do With Three Failures

Build the warning sign in first

Stub the step if it may not land

Give the risky box its best owner

Which Step Gets the Attention

	Intake	Parse	Score
Unknown	Yes	Yes	Yes
Blocks rest	No	Yes	No
Hard to fix	Yes	No	No

Write the Risk Down

Put it on the diagram

Name the warning sign

Say who is watching for it

Done With This Session When

The map runs input to output

It is an Artifact with a link

Every box names what it needs

The riskiest step has three failures

What You Leave With

1

Shared diagram

4

Kinds of need

3

Named failures

Next: Dividing the Work

Give every step one owner

Run workstreams in parallel

Put the board in GitHub

| Which box comes last?

Dividing the Work

Immersion Week SF · Day 1 · Session 1.3

50



Minutes

You leave with a board and owners

What This Session Covers

- 01** Inventory the team's real skills
- 02** Give every step a single owner
- 03** Run workstreams with Cowork
- 04** Put the board in GitHub

| Who owns which box?

From 1.2 to Now

A map

The path, drawn

vs

A board

The path, assigned

How We Get There

01

Skills

What can
we each do

02

Own

One name
per step

03

Split

Run them
in parallel

04

Board

Issues in
GitHub



01

Owners, Not Volunteers



Skills inventory

An honest list of what each person can do unaided.

What Counts as a Skill

Something you have shipped before

A tool you already have access to

A domain you actually know

Say the Quiet Part

What you have never done

What you would need help with

What you would rather not own

Skill to Step

Skill	Steps it fits
Data wrangling	Intake, parse
Prompt design	Score, summarise
Front end	Report, demo
Domain expert	Judging quality



Single owner

One named person accountable for a step being done.

Why Exactly One

Shared

Everyone assumes cover

Nobody watches the clock

Slips found late

Single

One person tracking

One person to ask

Slips found early

Assign Every Step

01 Riskiest first

It gets the best fit

02 Then the rest

Match skill to step

03 Name the gaps

Steps nobody can own

Volunteer or Best Fit

Volunteer

Picks what excites them

Learns on your deadline

Risky step goes last

Best fit

Picks what they can ship

Learns on the safe step

Risky step goes first

When Nobody Can Own a Step

Stub it and move on

Buy it with an existing tool

Cut it from the happy path



"One name per box, or no name at all."

- The rule that survives Day 2

Roles Beyond Writing Code

Owner

Holds the
problem

Tester

Breaks it on
purpose

Scribe

Keeps the
Project
current

Demo

Tells the
story Friday



02

Work in Parallel



Claude Cowork

Claude running multi-step work on files you choose.

Queue or Parallel

Queue

One step at a time

Everyone waits on one

Day 2 runs out

Parallel

Streams run at once

Waiting is designed out

Day 2 finishes

What Cowork Runs

On your files

Reads and edits what you pick

Multi-step

Finishes a task, not a reply

Alongside you

You work while it works

Set Up a Workstream

01 One step only

Scope it to a single box

02 Give the contract

Its input and output

03 Say done

How it proves it worked



A Workstream Brief

Step: Parse resumes to text

Input: PDF files in /intake

Output: one JSON per file
 {name, skills[], years}

Done when: 10 sample files
parse with no manual fixes.



The contract

The agreed shape passed between two steps.

Agree Contracts First

Name every field you pass

Decide it before building

Write it on the diagram

Step to Contract

Handoff

What passes

Intake to parse

File path plus id

Parse to score

JSON with skills

Score to report

Id plus a number

How Streams Collide

Two people edit one file

A contract changed silently

One stream waits on another

Keeping Streams Apart

One owner per file or folder

Change contracts by announcement

Stub what you are waiting on



03

Put It in GitHub

Three Things You Need

A repo

One place the code lives

Issues

One per task, with an owner

Done

How each task proves itself

Stand It Up Now

01 Create the repo

Name it for the problem

02 Add everyone

Write access, today

03 One issue per step

Straight from the map

An Issue, Filled In

Title: Parse resumes to text

Owner: @teammate

Input: PDF files in /intake

Output: one JSON per file

Done when: 10 samples parse
with no manual fixes.

Definition of done

A check someone else can run to confirm it works.

A Good Definition of Done

Observable, not felt

Checkable in a minute

Written before the work

Task to Done

Task	Done when
Parse files	10 samples, no fixes
Score match	Ranks a known set right
Build report	One click, no errors

Issue Hygiene

One task, one issue

Close it the moment it is done

Comment when you are blocked

Is This Task Ready?

	Parse	Score	Deploy
Has owner	Yes	Yes	Yes
Has contract	Yes	Yes	No
Has done	No	No	No

Ready for Tomorrow When

Every step is an issue

Every issue has one owner

Every issue says done

Contracts are written down

What You Leave With

1

Repo and board

1

Owner per step

0

Unowned tasks

Day 1, Behind You

A precise problem statement

A workflow map, shared

A board with owners

Next: Choosing the Stack

Match each step to a tool

Decide what to stub

Pick retrieval only if needed

| Can you start unaided?

Choosing the Stack

Immersion Week SF · Day 2 · Session 2.1

35



Minutes

You leave with a tool per step

What This Session Covers

- 01** Give every step a named tool
- 02** Decide what to stub, not build
- 03** RAG or just more context
- 04** Stand up the smallest layer

| What will you not build?

From 1.3 to Now

A board

Who does each step

vs

A stack

What each step runs on

How We Get There

01

Map

Reopen
yesterday's
diagram

02

Tool

Name one
per step

03

Stub

Mark what
stays fake

04

Retrieve

Only if you
must



"Boring tools, interesting problem."

- The Day 2 rule



01

A Tool for Every Step



Low-code

Assembling built pieces instead of writing them

The Rule for This Hour

Every box gets one tool name

No box may stay undecided

"We fake it" counts as decided

Write it on the board now

Where the Tool Names Come From

What your team already uses

What the workshop hands you

What has a free tier today

Not what you meant to learn

The Four Jobs

Intake

Getting the
input in

Think

The model
doing the
work

Store

Anything
you look up
later

Show

What the
user
actually
sees

A Worked Mapping

Step	Tool
Collect the input	A form or a paste box
Understand it	Claude, one prompt
Look something up	Claude Project files
Produce the output	Claude Artifact
Show a stranger	A hosted page

Signs You Picked Wrong

Nobody on the team has used it

Setup is measured in hours

It needs a key you don't have

The docs assume production

Build or Stub

Build it

It is the thing you claim

The demo touches it live

Nobody believes it faked

Stub it

It is plumbing, not proof

The demo never opens it

It works but is boring



Stub

A fake that is honest about being fake

How to Decide, Per Box

01 Ask what it proves

If nothing, it is plumbing

02 Ask who sees it

If nobody, stub it

03 Ask what it costs

Over an hour needs a reason

Always Safe to Stub

Login and user accounts

Payments of any kind

Email and notifications

Anything with a real database

Never Safe to Stub

The step your idea rests on

Whatever you claim is novel

The output a stranger reads

The path the demo walks

Three Tests Before You Commit

Under an hour

You can have it running before lunch

Someone knows it

One person on the team has shipped with it

Fails quietly

When it breaks, the demo still walks

Picking Between Two Options

	Familiar tool	Better tool
Setup	Minutes	Half a day
Ceiling	Low	High
Demo risk	Low	Unknown

Where the Time Actually Goes

80%

Wiring things together

Not the clever part

20%

The idea itself

What you are judged on



02

RAG or Just Context



RAG

Fetch the few relevant bits, then ask

| Does it all fit already?

Two Ways to Give Context

Paste it

Everything, every time

VS

Fetch it

Only what is relevant

What a Context Window Holds

A long policy document

A few hundred support tickets

A codebase of modest size

More than you probably have

The Decision, in Order

01 Measure the corpus

Actual size, not your guess

02 Try pasting it

If it fits, you are done

03 Only then, retrieve

And keep it as small as you can

Which One Do You Need

Just context

It fits in one window

The corpus barely changes

Every question sees it all

Real retrieval

Far more than fits

It changes constantly

Questions need a slice

What Each One Costs You

	Paste it	Retrieve it
Build time	Zero	Hours
Per question	More tokens	Fewer tokens
Breaks when	Corpus grows	Search misses
Debuggable	Easily	Less easily

The Costs Nobody Budgets

Chunking the documents well

Deciding how many to fetch

Debugging what it missed

Explaining it in the demo



Measure before you architect

```
$ wc -w docs/*.md  
41,204 total
```

~55k tokens. It fits.

No retrieval needed today.

You Are Over-Building If

You picked a vector database first

Your corpus is under 20 files

You cannot name a missed query

It is not tomorrow's demo step



"The cheapest thing that works."

- Objective two, in five words



03

The Smallest Layer



Claude Project

A workspace with your files always in context

Why It Is the Baseline

Standing it up takes minutes

No keys, no store, no code

Your files are simply there

Anyone can drive it

Stand One Up

01 Create the Project

Name it after the problem

02 Add your documents

The real ones, not samples

03 Write the instructions

Say how to answer, and how not to



Project instructions worth writing

Answer only **from** the attached files.

Quote the filename you used.

If the files do **not** cover it,
say so plainly. Do **not** guess.

What Goes in the Project

The documents you must cite

One example of a good answer

Your glossary and odd terms

Nothing you would not ship

The Baseline to Beat

Measure	Your target
Questions answered	More than the Project
Wrong answers	Fewer than the Project
Time to build	Under one afternoon
Teammate can run it	Yes

How to Know You Beat It

01 Write ten questions

Real ones, before you build

02 Run both, same questions

Project first, then yours

03 Count, do not vibe

Keep whichever wins

Three Layers, Cheapest First

Paste it in

Zero build; use it while the corpus is small

Claude Project

Minutes to stand up; the baseline to beat

Real retrieval

Only once you can name what the others missed

When to Graduate

The Project cannot hold it all

You can name what it missed

The corpus changes daily

You have the afternoon spare

Two Teams, Same Corpus

Built retrieval

Afternoon on chunking

Demo has a search step

Nobody else can run it

Used a Project

Twenty minutes total

Demo has an answer

Whole team can run it

Leaving This Session

Every box has a tool name

Every box is build or stub

Retrieval decided, not deferred

The baseline is standing up

Write This Down Now

Question	Your answer
Riskiest tool choice	...
Biggest stub	...
Retrieval: yes or no	...
What proves it works	...

| Can you start building?

Next: Building the Happy Path

One known-good input, end to end

No edge cases, no error handling

Claude Code drives each step

Commit the moment it works

The Happy Path

Immersion Week SF · Day 2 · Session 2.2

55



Minutes

One input reaches one output

What This Session Covers

- 01** Pick one known-good input
- 02** Build each step against it
- 03** Connect the steps end to end
- 04** Commit the moment it works

| What is your one input?

From 2.1 to Now

A stack

What each step runs on

vs

A path

One input, all the way

How We Get There

01

Input

Choose
one, freeze
it

02

Build

One step at
a time

03

Chain

Remove
every hand

04

Commit

The second
it works

The Next 55 Minutes

10:00

Freeze it

One file,
committed,
shared

10:10

Build steps

One box at
a time,
verified

10:35

Connect

Run it from
a cold start

10:50

Commit

Tag the
version that
works



"Make it work. Only then, anything else."

- The Day 2 rule



01

One Known-Good Input



Happy path

The one route where nothing goes wrong

What It Is Not

Not the average case

Not most of your inputs

Not robust to anything

Not what you ship later

Choosing the Input

01 Take a real one

From the actual problem, not invented

02 Check it is boring

Nothing weird, nothing missing

03 Freeze it in the repo

Same bytes for everyone, all day

What Makes One Good

Small enough to read fully

Real enough to be convincing

Boring enough to always work

Fixed, so results compare



Freeze it where everyone can see

```
$ mkdir -p fixtures
```

```
$ cp ~/Downloads/report.pdf fixtures/
```

```
$ git add fixtures/report.pdf
```

```
$ git commit -m "freeze demo input"
```

Today and Not Today

Today

One input, all the way

The steps from your map

Whatever is ugliest

Not today

Edge cases

Error handling

A second input

The Ban List for This Hour

No try or catch anywhere

No configuration files

No second input to test

No renaming or refactoring



Claude Code

An agent that edits and runs your repo

Driving It, Per Step

01 Name the one step

From your map, by its box

02 Give it the contract

What comes in, what goes out

03 Point at the input

The frozen file, by path

Give It What It Needs

The map, as a file in the repo

The contracts between steps

The frozen input, by path

The command that proves it ran

Two Ways to Ask

Vague

"Parse the document"

"Make it handle PDFs"

"Clean up the output"

Specific

"Read fixtures/report.pdf"

"Return title and date"

"Write it to out.json"



A request worth making

Read fixtures/report.pdf.

Pull out title, author, date.

Write them to out/step1.json.

Run it **and** show me the file.

Three Rules for Driving

One step per ask

Never two boxes of the map at once

Make it run

Ask for the output, not a description

Reset, do not nudge

If two tries fail, restate the whole ask

When It Goes Sideways

It invented a file — name paths

It over-built — say what to skip

It cannot run — give it the command

It loops — restate from scratch

Check Each Step Before Moving On

You ran it yourself, just now

The output file exists

You read the output

It matches the contract



02

Connect the Steps

| Where do hands still touch?



Manual step

Anything a person does between two steps

Manual Steps That Hide

Copying output into the next box

Renaming a file by hand

Pasting into a chat window

Knowing to run them in order

Two Kinds of Demo

Narrated

And now I open this

vs

Run once

One command, done

Connecting, in Order

01 Write the contracts down

What each step leaves behind

02 Chain two, then three

Add one box at a time

03 Run it from cold

Delete the outputs and start over

Write Down the Seams

Step	Leaves behind
------	---------------

Ingest	out/raw.json
--------	--------------

Extract	out/fields.json
---------	-----------------

Reason	out/answer.md
--------	---------------

Render	out/index.html
--------	----------------



The whole path, one command

```
$ rm -rf out/ && mkdir out
```

```
$ python run.py fixtures/report.pdf
```

```
... 4 steps ok
```

```
$ open out/index.html
```

The Seam Is Where It Breaks

One step writes, the next cannot read

A field is named two ways

One step assumes a list

A path is right on one machine

Debugging a Seam

01 Open both files

What was written, what was expected

02 Fix the later step

Adapt downstream, not upstream

03 Rerun from cold

Never trust a partial run

Connect Early or Late

	Connect early	Connect late
Surprises	Now, small	Late, large
Rework	Cheap	Expensive
Demo risk	Low	High

Four Smells

Only Ana

One person
can run it

Warm only

Fails from a
clean start

In order

You must
know the
sequence

One box

A step still
done by
hand

If You Are Behind

01 Cut a box, not a corner

Stub a whole step, keep the path whole

02 Shrink the input

One page instead of the whole report

03 Never drop the commit

A short path committed beats a long one lost

Done Means All Four

One command, start to finish

Runs from an empty output dir

Someone else ran it

The output is the real thing



"If hands touch it, it is not connected."

- Objective two, restated



03

Commit It Immediately



Works

It ran end to end, cold, in front of you

Why Right Now

This is the version that works

The next change may break it

You will not remember what changed

A demo needs a known-good commit

The Commit Ritual

01 Stop typing

Hands off the keyboard, it works

02 Commit and push

Message says what now works

03 Then keep going

Anything new lands on top of it



Sixty seconds of insurance

```
$ git add run.py steps/ fixtures/
```

```
$ git commit -m "happy path runs cold"
```

```
$ git push -u origin main
```

```
$ git tag happy-path && git push --tags
```

What the repo looks like now

```
happy-path/  
  fixtures/  
    report.pdf  
  steps/  
    extract.py  
    render.py  
  run.py  
  README.md
```


What the Message Should Say

What input it accepts

What output it produces

What is still stubbed

Not how clever it is

Before and After the Commit

Before committing

Nothing else at all

No tidying, no renaming

No one more small fix

After committing

Everything you want

Break it freely

You can always get back

What the Commit Buys You

A place to stand

Every later change can be undone to here

A shared build

Teammates start from what already works

A safe demo

A known-good tag if tomorrow goes wrong

Then, and Only Then

Add a second input

Handle the obvious failure

Make the output look right

Start on the demo wrapper

Write This Down Now

Question	Your answer
The frozen input	...
The one command	...
The known-good tag	...
Still done by hand	...

Leaving This Session

One input, frozen in the repo

One command runs it all

It works from a cold start

It is committed and pushed

| Could a stranger run it?

Next: Making It Demoable

A front end a stranger can drive

A public URL, not your laptop

A three-minute demo script

An honest list of what is faked

Interface Aesthetics

Ngo, Teo & Byrne · Information Sciences (2003)



Is beauty computable?



Interface aesthetics

Layout quality, written as numbers.

14

Aesthetic measures

Closed-form functions of element bounding boxes

The core claim

Layout quality is not purely subjective.

Balance and spacing become numbers.

Each number is a formula over boxes.

Measured, optimised, checked vs people.



01

The fourteen measures

The measures (1–5)

Balance — even visual weight

Equilibrium — mass near the centre

Symmetry — mirrored about the axes

Sequence — a clear reading flow

Cohesion — aspect ratios agree



Bounding box

Every measure reads position and size.

The measures (6–10)

Unity — parts read as one

Proportion — pleasing ratios

Simplicity — few alignment points

Density — comfortable coverage

Regularity — elements snap to a grid

The measures (11–14)

Economy — few sizes, reused

Homogeneity — evenly distributed

Rhythm — regular, repeating change

Order & complexity — structure wins

regularity — Ngo (2003)

Ngo regularity (alignment term)

$$\text{regularity} = 1 - (n_{\text{vap}} + n_{\text{hap}}) / (2n)$$

a tidy 5-bullet column:

$n = 5$, $n_{\text{vap}} = 1$, $n_{\text{hap}} = 5$

$$\text{regularity} = 1 - (1 + 5) / 10 \quad \# 0.40$$

slidekit: left edges only

$$\text{regularity} = 1 - 1 / 5 \quad \# 0.80$$

Same clean column, two scores

0.40

Ngo x + y regularity

vs

0.80

Left-edge alignment

Why slidekit fits this paper

Every element already has an exact rect.

Measures are arithmetic over geometry.

Deterministic — same deck, same score.

Zero rendering, zero vision.



02

From paper to slidekit

What slidekit adopted

Declined

Regularity

Rhythm

Proportion

Economy

Spacing-regularity

Adopted

Balance

Density

Adoption decisions

Measure	Verdict	Reason
Balance	Adopted	Axis moment
Density	Adopted	Coverage
Regularity	Declined	Halves cols
Proportion	Declined	Ratio noise

A measured-aesthetics lineage

1933

Birkhoff

A beauty
ratio

2003

Ngo

Fourteen
measures

2004

Lok

Visual
balance

2026

Zhang

Moment
model

Honest scope

The journal paper is paywalled.

Public copies reproduce only alignment.

Constants are not copied blindly.

No universal beauty equation is claimed.

From measure to code

01 Read geometry

Take each element's rect.

02 Define a formula

Write an open, closed-form score.

03 Verify & ship

Check the source, then report.



"Make design quality auditable."

- Ngo, Teo & Byrne (2003)