

# Victor Palacios

Director of Data Science and AI Partnerships · USF

**| How do you end up here?**

# Still Collecting Degrees

**2011**

**Berkeley**

BA in  
Japanese  
language  
and  
literature

**2015**

**Nagoya**

MS in  
Information  
Science, in  
Japan

**2021**

**USF**

MS in Data  
Science

**2026**

**Quantic**

MS in  
Software  
Eng, under  
way

# What I Will Cover

---

- 01** Research, from brains to health
- 02** The path, and its detours
- 03** Writing things down in public
- 04** Partnerships, and one big repo



**01**

**Research**



# Deception detection

Finding what the brain does when someone lies

# Nagoya, 2014

Are action lies unlike speech lies?

Measured with near-infrared light

Forty sentences, three sessions

Published at JCSS in Japan

# How the Study Ran

## **01 Learning**

You gesture or speak twenty sentences

## **02 Recognition**

You recall which ones you did which way

## **03 Lying**

You lie about twenty, tell the truth about twenty



**"Different lies, different regions."**

- Palacios, Saito et al., JCSS 2014

# Then Healthcare, 2021

A sensor reads blood sugar all day

Can it tell when you ate?

Healthy participants, not patients

First author, with Woodbridge and Fry

# What It Was Really For

**Asked**

vs

**Measured**

Did you eat when told?

The sensor knew

# The Number That Mattered

**60%**

Baseline

Where we started

**89%**

After the work

Same data, better target

# ICD-10 Coding, 2024

Clinical notes into diagnosis codes

Small open models, not large ones

Over 100,000 notes from MIMIC-IV

Swiss and USF collaboration

# What We Benchmarked

| Model    | Size            |
|----------|-----------------|
| Falcon   | 1B parameters   |
| Phi-3.5  | 3.5B parameters |
| Meditron | 7B parameters   |

# What Surprised Us

The 1B model held its own

It was not a medical model

LoRA kept quality, cut compute

Small can be good enough

# Three Papers, One Habit

| Year | Question                   |
|------|----------------------------|
| 2014 | What does a lie look like? |
| 2021 | Did you really eat?        |
| 2024 | Can a small model code?    |



**02**

# **The Path**

# Where I Started

**Japanese**

Language, literature

vs

**Now**

Data science and AI

# What Japan Taught Me

Language is structured data

Research is mostly patience

Being the outsider is normal

Fluency takes years, not weeks

# Industry, in Between

Halo Solutions, working with the CEO

Found overdue patients in 10,000 PDFs

KNIME, teaching six courses

Fraud detection, to a big audience

# Two Numbers From Industry

**\$13M**

Revenue identified

Overdue patients found

**90k**

Subscriber audience

For the fraud talk

# What I Do Now

Match students to practicum teams

Mentor those teams for nine months

Teach the SQL certificate course

Career advising, all year

# Still a Student Myself

A software engineering master's

Started October 2025

While doing this job

The learning does not stop



**03**

## **Writing It Down**

# mAlstory

Fine-tuned Gemma 2B for stories

An AI storyteller for children

Built through KaggleX mentorship

Code published with the article

# The AWS Certificate

Cloud Practitioner, passed first go

Zero cloud experience going in

Only free study materials

Wrote up exactly how

# The Numbers From That Post

**838**

Score

700 is a pass

**32**

Hours

Across one month

**\$0**

Spent

Free materials only



**"The paywall is usually optional."**

- A recurring theme



**04**

# **Partnerships**

# The Practicum

Sixteen hours a week

Roughly nine months

Every student, guaranteed

Real problems, real clients

# The Scale of It

**120+**

Companies

Relationships I hold

**150+**

Students

Taught and mentored

**5.67**

Of six

My teaching score

# Who We Work With

## Finance

BlackRock

## Cloud

AWS and  
Atlassian

## Health

UCSF,  
Stanford

## Mission

The Gates  
Foundation

# Where Our Graduates Land

**78%**

Take offers in California

Class of 2025

**\$150k**

Median base salary

California offers

# What Gets Someone Hired

Work a stranger can look at

One project explained clearly

Curiosity that survives boredom

Not a perfect transcript



**05**

# **Ninth Circuit**

# The Problem

Thousands of opinions a year

Only a fraction are asylum cases

Reading them all is impossible

So the research does not happen

# How the Pipeline Works

**01**

**Fetch**

Scrape opinions

**02**

**Classify**

Asylum or not?

**03**

**Extract**

Pull the features

# What It Has Built

**5,779**

Opinions

Growing daily

**70+**

Features each

With a source quote

**26%**

Are asylum

The rest filtered

# The Trick That Makes It Free

Classify cheaply before extracting

Throw away three quarters first

Roughly four times cheaper

Free-tier models throughout

# What It Runs On

Llama 3.3 on a free API tier

Postgres, and a vector index

A daily scheduled job

Total ongoing cost: nothing

# Why It Matters

This work was always hand-coded

Now it is searchable

Every claim has its quote

Anyone can read the code

# What I Actually Believe

**Paths are not straight**

---

I studied literature first

Nothing was wasted

The detour is the edge

**The tools are free**

---

A certificate, for nothing

A legal database, free

The barrier is starting

# If You Take One Action

Build one small thing publicly

Write down how you did it

Use the free tier for all of it

Then show someone

**| What will you build?**